

Psamathe: A Flow-Based DSL for Safe Blockchain Assets

Reed Oei¹, Michael Coblenz², and Jonathan Aldrich³

¹ University of Illinois, Urbana, IL, USA
`reedoei2@illinois.edu`

² University of Maryland, College Park, MD, USA
`mcoblenz@umd.edu`

³ Carnegie Mellon University, Pittsburgh, PA, USA
`jonathan.aldrich@cs.cmu.edu`

1 Introduction

Blockchains are increasingly used as platforms for applications called *smart contracts* [14], which automatically manage transactions in an mutually agreed-upon way. Commonly proposed and implemented applications include supply chain management, healthcare, voting, crowdfunding, auctions, and more [9,8,7]. Smart contracts often manage *digital assets*, such as cryptocurrencies, or, depending on the application, bids in an auction, votes in an election, and so on. These contracts cannot be patched after deployment, even if security vulnerabilities are discovered. Some estimates suggest that as many as 46% of smart contracts may have vulnerabilities [10].

Psamathe (/sɑmɑθi/) is a new programming language we are designing around *flows*, which are a new abstraction representing an atomic transfer operation. Together with features such as *modifiers*, flows provide a **concise** way to write contracts that **safely** manage assets (see Section 2). Solidity, the most commonly-used smart contract language on the Ethereum blockchain [1], does not provide analogous support for managing assets. A formalization of Psamathe is in progress [2], with an *executable semantics* implemented in the \mathbb{K} -framework [11], which is already capable of running the example shown in Figure 1.

Other newly-proposed blockchain languages include Flint, Move, Nomos, Obsidian, and Scilla [12,4,6,5,13]. Scilla and Move are intermediate-level languages, whereas Psamathe is intended to be a high-level language. Obsidian, Move, Nomos, and Flint use linear or affine types to manage assets; Psamathe uses *type quantities*, which provide the benefits of linear types, but allow a more precise analysis of the flow of values in a program.

2 Language

A Psamathe program is made of *transformers* and *type declarations*. Transformers contain *flows* describing the how values are transferred between variables. Type declarations provide a way to mark values with *modifiers*. Figure 1 shows a

simple contract declaring a type and a transformer, which implements the core of ERC-20’s `transfer` function. ERC-20 is a standard providing a bare-bones interface for token contracts managing *fungible* tokens—all tokens of an ERC-20 compliant contract are interchangeable (like most currencies), so it is only important how many tokens are owned by an entity, not **which** tokens.

```

1 type Token is fungible asset uint256
2 transformer transfer(balances : any map ! address => any Token,
3                     dst : ! address, amount : any uint256) {
4   balances[msg.sender] --[ amount ]-> balances[dst]
5 }
```

Fig. 1: A Psamathe contract with a simple transfer function, which transfers `amount` tokens from the sender’s account to the destination account. It is implemented with a single flow, which automatically checks all the preconditions to ensure the transfer is valid.

Psamathe is built around the concept of a flow. Using the more declarative, *flow-based* approach provides the following advantages:

- **Static safety guarantees:** Each flow is guaranteed to preserve the total amount of assets (except for flows that explicitly consume or allocate assets). The total amount of a nonconsumable asset never decreases. Each asset has exactly one reference to it, either via a variable in the current environment, or in a table/record. The `immutable` modifier prevents values from changing.
- **Dynamic safety guarantees:** Psamathe automatically inserts dynamic checks of a flow’s validity; e.g., a flow of money would fail if there is not enough money in the source, or if there is too much in the destination (e.g., due to overflow). The `unique` modifier, which restrict values to never be created more than once, is also checked dynamically.
- **Data-flow tracking:** We hypothesize that flows provide a clearer way of specifying how resources flow in the code itself, which may be less apparent using other approaches, especially in complicated contracts. Additionally, developers must explicitly mark when assets are *consumed*, and only assets marked as `consumable` may be consumed.
- **Error messages:** When a flow fails, the Psamathe runtime provides automatic, descriptive error messages, such as

```
Cannot flow <amount> Token from account[<src>] to account[<dst>]:
source only has <balance> Token.
```

Flows enable such messages by encoding information into the source code.

Each variable and function parameter has a *type quantity*, approximating the number of values, which is one of: `empty`, `any`, `!`, `nonempty`, or `every` (“!” means “exactly one”). Only `empty` asset variables may be dropped. Type quantities are inferred if omitted; every type quantity in Figure 1 can be omitted.

Modifiers can be used to place constraints on how values are managed: they are `asset`, `consumable`, `fungible`, `unique`, and `immutable`. An `asset` is a value that

must not be reused or accidentally lost, such as money. A `consumable` value is an `asset` that it may be appropriate to dispose of, via the `consume` construct, documenting that the disposal is intentional. For example, while bids should not be lost **during** an auction, it is safe to dispose of them after the auction ends. A `fungible` value can be **merged**, and it is **not unique**. The modifiers `unique` and `immutable` provide the safety guarantees mentioned above.

One could try automatically inserting dynamic checks in a language like Solidity, but in many cases it would require additional annotations. Such a system would essentially reimplement flows, providing some benefits of Psamathe, but not the same static guarantees.

Comparison with ERC-20 in Solidity Each ERC-20 contract manages the “bank accounts” for its own tokens, keeping track of how many tokens each account has; accounts are identified by addresses. We compare the Psamathe implementation in Figure 1 to Figure 2, which shows a Solidity implementation of the same function. In this case, the sender’s balance must be at least as large as `amount`, and the destination’s balance must not overflow when it receives the tokens. Psamathe automatically inserts code checking these two conditions, ensuring the checks are not forgotten.

```

1 contract ERC20 {
2   mapping (address => uint256) balances;
3   function transfer(address dst, uint256 amount) public {
4     require(amount <= balances[msg.sender]);
5     balances[msg.sender] = balances[msg.sender].sub(amount);
6     balances[dst] = balances[dst].add(amount);
7   }
8 }
```

Fig. 2: An implementation of ERC-20’s `transfer` function in Solidity from one of the reference implementations [3]. All preconditions are checked manually. Note that we must include the `SafeMath` library (not shown) to use the `add` and `sub` functions, which check for underflow/overflow.

3 Conclusion and Future Work

We have presented the Psamathe language for writing safer smart contracts. Psamathe uses the new flow abstraction, assets, and type quantities to provide its safety guarantees. We have shown an example smart contract in both Psamathe and Solidity, showing that Psamathe is capable of expressing common smart contract functionality in a concise manner, while retaining key safety properties.

In the future, we plan to implement the Psamathe language, and prove its safety properties. We also hope to study the benefits and costs of the language via case studies, performance evaluation, and the application of flows to other domains. Finally, we would also like to conduct a user study to evaluate the usability of the flow abstraction and the design of the language, and to compare

it to Solidity, which we hypothesize will show that developers write contracts with fewer asset management errors in Psamathe than in Solidity.

References

1. Ethereum for developers (2020), <https://ethereum.org/en/developers/>
2. Psamathe (Aug 2020), <https://github.com/ReedOei/Psamathe>
3. Tokens (Aug 2020), <https://github.com/ConsenSys/Tokens>
4. Blackshear, S., Cheng, E., Dill, D.L., Gao, V., Maurer, B., Nowacki, T., Pott, A., Qadeer, S., Rain, D.R., Sezer, S., et al.: Move: A language with programmable resources (2019)
5. Coblenz, M., Oei, R., Etzel, T., Koronkevich, P., Baker, M., Bloem, Y., Myers, B.A., Sunshine, J., Aldrich, J.: Obsidian: Typestate and assets for safer blockchain programming (2019)
6. Das, A., Balzer, S., Hoffmann, J., Pfenning, F., Santurkar, I.: Resource-aware session types for digital contracts. arXiv preprint arXiv:1902.06056 (2019)
7. Elsdén, C., Manohar, A., Briggs, J., Harding, M., Speed, C., Vines, J.: Making sense of blockchain applications: A typology for hci. In: CHI Conference on Human Factors in Computing Systems. pp. 1–14. CHI '18 (2018). <https://doi.org/10.1145/3173574.3174032>, <http://doi.acm.org/10.1145/3173574.3174032>
8. Harvard Business Review: The potential for blockchain to transform electronic health records (2017), <https://hbr.org/2017/03/the-potential-for-blockchain-to-transform-electronic-health-records>
9. IBM: Blockchain for supply chain (2019), <https://www.ibm.com/blockchain/supply-chain/>
10. Luu, L., Chu, D.H., Olickel, H., Saxena, P., Hobor, A.: Making smart contracts smarter. In: Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security. p. 254–269. CCS '16, Association for Computing Machinery, New York, NY, USA (2016). <https://doi.org/10.1145/2976749.2978309>, <https://doi.org/10.1145/2976749.2978309>
11. Roşu, G., Şerbănuţă, T.F.: An overview of the K semantic framework. *Journal of Logic and Algebraic Programming* **79**(6), 397–434 (2010). <https://doi.org/10.1016/j.jlap.2010.03.012>
12. Schrans, F., Eisenbach, S., Drossopoulou, S.: Writing safe smart contracts in flint. In: Conference Companion of the 2nd International Conference on Art, Science, and Engineering of Programming. pp. 218–219 (2018)
13. Sergey, I., Nagaraj, V., Johannsen, J., Kumar, A., Trunov, A., Hao, K.C.G.: Safer smart contract programming with scilla. *Proc. ACM Program. Lang.* **3**(OOPSLA) (Oct 2019). <https://doi.org/10.1145/3360611>, <https://doi.org/10.1145/3360611>
14. Szabo, N.: Formalizing and securing relationships on public networks. *First Monday* **2**(9) (1997). <https://doi.org/https://doi.org/10.5210/fm.v2i9.548>